

## ISLR Chapter 4

### Exercise 1

Using a little bit of algebra, use (4.2) to achieve (4.3). In other words, the logistic function representation and logit representation for the logistic regression model are equivalent. Where:

$$(4.2) \quad p(X) = \frac{e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}}$$

$$(4.3) \quad \frac{p(X)}{1 - p(X)} = e^{\beta_0 + \beta_1 X}$$

$$\text{So, } \frac{p(X)}{1 - p(X)} = \frac{\frac{e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}}}{1 - \frac{e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}}} = \frac{\frac{e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}}}{\frac{1 + e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}} - \frac{e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}}} = \frac{\frac{e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}}}{\frac{1}{1 + e^{\beta_0 + \beta_1 X}}} = e^{\beta_0 + \beta_1 X}$$

### Exercise 6

Suppose we collect data for a group of students in a statistics class with variables  $X_1$  = hours studied,  $X_2$  = undergrad GPA, and  $Y$  = receive an A. We fit a logistic regression and produce estimated coefficients:

$$\hat{\beta}_0 = -6 \quad \hat{\beta}_1 = 0.05 \quad \hat{\beta}_2 = 1$$

6. a) Estimate the probability that a student who studies for 40 h and has an undergrad GPA of 3.5 gets an A in the class.

Remember from the previous exercise that:

$$p(X) = \frac{\exp(\beta_0 + \beta_1 X_1 + \beta_2 X_2)}{1 + \exp(\beta_0 + \beta_1 X_1 + \beta_2 X_2)}$$

Estimating the probability for  $X = [40hours, 3.5GPA]$  yields:

$$p(X) = \frac{\exp(-6 + 0.05X_1 + X_2)}{1 + \exp(-6 + 0.05X_1 + X_2)} = \frac{\exp(-6 + 0.05 * 40 + 3.5)}{1 + \exp(-6 + 0.05 * 40 + 3.5)} = \frac{\exp(-0.5)}{1 + \exp(-0.5)} = 37.75\%$$

**6. b) How many hours would the student in part (a) need to study to have a 50% chance of getting an A in the class?**

Estimating  $X_1$  where  $X = [X_1 hours, 3.5GPA]$  such that  $p(X) = 0.5$  yields:

$$0.50 = \frac{\exp(-6 + 0.05X_1 + 3.5)}{1 + \exp(-6 + 0.05X_1 + 3.5)} \Leftrightarrow 0.50(1 + \exp(-2.5 + 0.05X_1)) = \exp(-2.5 + 0.05X_1)$$

$$\Leftrightarrow 0.50 + 0.50 \exp(-2.5 + 0.05X_1) = \exp(-2.5 + 0.05X_1) \Leftrightarrow 0.50 = 0.50 \exp(-2.5 + 0.05X_1)$$

$$\Leftrightarrow \log(1) = -2.5 + 0.05X_1 \Leftrightarrow X_1 = 2.5/0.05 = 50hours$$

## Exercise 10

This question should be answered using the Weekly data set, which is part of the ISLR package. This data is similar in nature to the Smarket data from this chapter's lab, except that it contains 1,089 weekly returns for 21 years, from the beginning of 1990 to the end of 2010.

10.a) Produce some numerical and graphical summaries of the Weekly data. Do there appear to be any patterns?

```

In [1]: #No warnings for print version
options(warn=-1)

#First, load the necessary R-packages:
#install.packages("ISLR")
#install.packages("MASS")
#install.packages("class")

#call the packages you need
library("ISLR")

#Take a first look at the data - this time no need to clean the data
fix(Weekly) # allows to edit the data
attach(Weekly) # eliminates the need of refering to a variable like Weekly$variable.

# Use summary function to produce a numerical summary for each variable
summary(Weekly)

```

Year	Lag1	Lag2	Lag3
Min. :1990	Min. : -18.1950	Min. : -18.1950	Min. : -18.1950
1st Qu.:1995	1st Qu.: -1.1540	1st Qu.: -1.1540	1st Qu.: -1.1580
Median :2000	Median : 0.2410	Median : 0.2410	Median : 0.2410
Mean :2000	Mean : 0.1506	Mean : 0.1511	Mean : 0.1472
3rd Qu.:2005	3rd Qu.: 1.4050	3rd Qu.: 1.4090	3rd Qu.: 1.4090
Max. :2010	Max. : 12.0260	Max. : 12.0260	Max. : 12.0260

Lag4	Lag5	Volume	Today
Min. : -18.1950	Min. : -18.1950	Min. : 0.08747	Min. : -18.1950
1st Qu.: -1.1580	1st Qu.: -1.1660	1st Qu.: 0.33202	1st Qu.: -1.1540
Median : 0.2380	Median : 0.2340	Median : 1.00268	Median : 0.2410
Mean : 0.1458	Mean : 0.1399	Mean : 1.57462	Mean : 0.1499
3rd Qu.: 1.4090	3rd Qu.: 1.4050	3rd Qu.: 2.05373	3rd Qu.: 1.4050
Max. : 12.0260	Max. : 12.0260	Max. : 9.32821	Max. : 12.0260

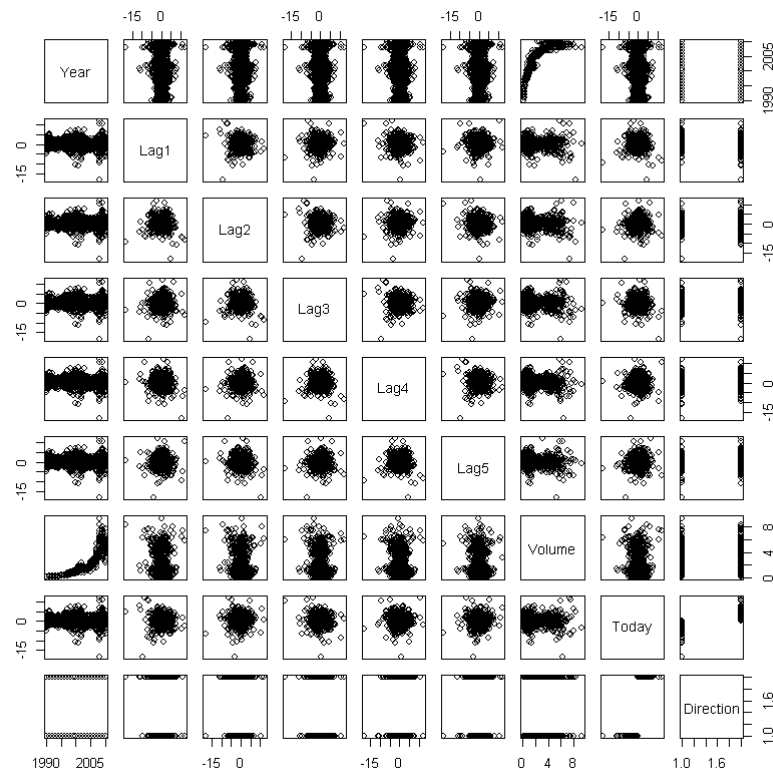
Direction  
 Down:484  
 Up :605

```
In [2]: # Use cor function to produce a table of correlations for all variables (excluding Direction non-numerical variable)
cor(Weekly[, -9])
```

A matrix: 8 × 8 of type dbl

	Year	Lag1	Lag2	Lag3	Lag4	Lag5	Volume	Today
Year	1.00000000	-0.032289274	-0.03339001	-0.03000649	-0.031127923	-0.030519101	0.84194162	-0.032459894
Lag1	-0.03228927	1.000000000	-0.07485305	0.05863568	-0.071273876	-0.008183096	-0.06495131	-0.075031842
Lag2	-0.03339001	-0.074853051	1.00000000	-0.07572091	0.058381535	-0.072499482	-0.08551314	0.059166717
Lag3	-0.03000649	0.058635682	-0.07572091	1.00000000	-0.075395865	0.060657175	-0.06928771	-0.071243639
Lag4	-0.03112792	-0.071273876	0.05838153	-0.07539587	1.000000000	-0.075675027	-0.06107462	-0.007825873
Lag5	-0.03051910	-0.008183096	-0.07249948	0.06065717	-0.075675027	1.000000000	-0.05851741	0.011012698
Volume	0.84194162	-0.064951313	-0.08551314	-0.06928771	-0.061074617	-0.058517414	1.00000000	-0.033077783
Today	-0.03245989	-0.075031842	0.05916672	-0.07124364	-0.007825873	0.011012698	-0.03307778	1.000000000

```
In [3]: # Use pairs function to produce a graphical summary
pairs(Weekly)
```



**Answer:** Yes, it appears that Year and Volume have a non-linear and positive relation.

**10. b) Use the full data set to perform a logistic regression with Direction as the response and the five lag variables plus Volume as predictors. Use the summary function to print the results. Do any of the predictors appear to be statistically significant? If so, which ones?**

```
In [4]: # Estimate a generalized linear regression model where the third input family is a description of the error distribution
# and link function to be used in the model, supplied as the result of a call to a family function - here use binomial.
# Why binomial? Because our independent variable Direction takes two values.

glm.fit = glm(Direction~Lag1+Lag2+Lag3+Lag4+Lag5+Volume,
               data=Weekly,
               family=binomial)

# Use summary function to print the results
summary(glm.fit)
```

Call:

```
glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
     Volume, family = binomial, data = Weekly)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.6949	-1.2565	0.9913	1.0849	1.4579

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	0.26686	0.08593	3.106	0.0019 **
Lag1	-0.04127	0.02641	-1.563	0.1181
Lag2	0.05844	0.02686	2.175	0.0296 *
Lag3	-0.01606	0.02666	-0.602	0.5469
Lag4	-0.02779	0.02646	-1.050	0.2937
Lag5	-0.01447	0.02638	-0.549	0.5833
Volume	-0.02274	0.03690	-0.616	0.5377

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1496.2 on 1088 degrees of freedom  
 Residual deviance: 1486.4 on 1082 degrees of freedom  
 AIC: 1500.4

Number of Fisher Scoring iterations: 4

**Answer:** the predictor Lag 2 appears to have some statistical significance with a p-value smaller than 3%.

**10. c) Compute the confusion matrix and overall fraction of correct predictions. Explain what the confusion matrix is telling you about the types of mistakes made by logistic regression.**

```
In [5]: # use predict function on results of previous regression in 10.b)
glm.probs = predict(glm.fit, type="response")

#create dataframe that is a vector always taking string "Down" same size as #obs
glm.pred = rep("Down", length(glm.probs))

#substitute for "Up" whenever the estimated probability is above 0.5
glm.pred[glm.probs>.5] = "Up"

#construct a summary table with the prediction against the actual values of the variable Direction
table(glm.pred, Direction)
```

```
      Direction
glm.pred Down  Up
Down      54  48
Up       430 557
```

**Answer:**

Percentage of correct predictions:  $(54+557)/(54+557+48+430) = 56.1\%$ .

During weeks when the market goes up, the logistic regression is right about  $557/(557+48) = 92.1\%$  of the time.

During weeks when the market goes down, the logistic regression is right about  $54/(430+54) = 11.2\%$  of the time.

**10. d) Now fit the logistic regression model using a training data period from 1990 to 2008, with Lag2 as the only predictor. Compute the confusion matrix and the overall fraction of correct predictions for the held out data (that is, the data from 2009 and 2010).**



```

In [6]: # generate condition for our training data
train = (Year < 2009)

# create dataframe for the Weekly data from 2009 and 2010 (usage of ! to define the "opposite")
Weekly.0910 = Weekly[!train,]

# run regression on the training data subset
glm.fit = glm(Direction~Lag2,
               data=Weekly,
               family=binomial,
               subset=train)

# create dataframe
glm.probs = predict(glm.fit, Weekly.0910,type="response")

# fill with our predictions
glm.pred = rep("Down", length(glm.probs))
glm.pred[glm.probs>.5] = "Up"

# construct confusion table using only the subset data
Direction.0910 = Direction[!train]
table(glm.pred, Direction.0910)

# overall fraction of correct predictions
mean(glm.pred == Direction.0910)

```

```

      Direction.0910
glm.pred Down Up
Down      9   5
Up       34  56

0.625

```

**10. e) Repeat (d) using LDA.**

```
In [7]: #call the packages you need
library("MASS")

# same approach as before but now using LDA method
lda.fit = lda(Direction ~ Lag2, data=Weekly, subset=train)
lda.pred = predict(lda.fit, Weekly.0910)
table(lda.pred$class, Direction.0910)
mean(lda.pred$class == Direction.0910)
```

```

      Direction.0910
      Down Up
Down      9  5
Up       34 56

0.625
```

#### 10. f) Repeat (d) using QDA.

```
In [8]: # same approach as before but now using QDA method
qda.fit = qda(Direction~Lag2, data=Weekly, subset=train)
qda.class = predict(qda.fit, Weekly.0910)$class
table(qda.class, Direction.0910)
mean(qda.class == Direction.0910)
```

```

      Direction.0910
qda.class Down Up
      Down      0  0
      Up       43 61

0.586538461538462
```

#### 10. g) Repeat (d) using KNN with K = 1.

```
In [9]: #call the packages you need
library("class")

# same approach as before but now using KNN method with K=1
train.X = as.matrix(Lag2[train])
test.X = as.matrix(Lag2[!train])
train.Direction = Direction[train]

#set seed to get the same results - if several observations are tied as nearest neighbors, then R will randomly break the tie.
#In order to break the tie the same way each time, you set the seed so you can reproduce the results exactly
set.seed(1)

#KNN prediction uses a different function
# documentationn here: https://www.rdocumentation.org/packages/class/versions/7.3-17/topics/knn
knn.pred = knn(train.X, test.X, train.Direction, k=1)
table(knn.pred, Direction.0910)
mean(knn.pred == Direction.0910)
```

```
      Direction.0910
knn.pred Down Up
Down      21 30
Up        22 31
```

```
0.5
```

**10. h) Which of these methods appears to provide the best results on this data?**

Logistic regression and LDA have the smallest test error rates.

**10. i) Experiment with different combinations of predictors, including possible transformations and interactions, for each of the methods. Report the variables, method, and associated confusion matrix that appears to provide the best results on the held out data. Note that you should also experiment with values for K in the KNN classifier.**

```
In [10]: # Logistic regression with Lag2:Lag1
glm.fit = glm(Direction~Lag2:Lag1, data=Weekly, family=binomial, subset=train)
glm.probs = predict(glm.fit, Weekly.0910, type="response")
glm.pred = rep("Down", length(glm.probs))
glm.pred[glm.probs>.5] = "Up"
Direction.0910 = Direction[!train]
table(glm.pred, Direction.0910)
mean(glm.pred == Direction.0910)
```

```

      Direction.0910
glm.pred Down Up
Down      1   1
Up       42  60
```

```
0.586538461538462
```

```
In [11]: # LDA with Lag2 interaction with Lag1
lda.fit = lda(Direction ~ Lag2:Lag1, data=Weekly, subset=train)
lda.pred = predict(lda.fit, Weekly.0910)
mean(lda.pred$class == Direction.0910)
```

```
0.576923076923077
```

```
In [12]: # QDA with sqrt(abs(Lag2))
qda.fit = qda(Direction~Lag2+sqrt(abs(Lag2)), data=Weekly, subset=train)
qda.class = predict(qda.fit, Weekly.0910)$class
table(qda.class, Direction.0910)
mean(qda.class == Direction.0910)
```

```

      Direction.0910
qda.class Down Up
Down      12  13
Up       31  48
```

```
0.576923076923077
```

```
In [13]: # KNN k =10, as before KNN uses a different command
set.seed(1)
knn.pred = knn(train.X, test.X, train.Direction, k=10)
table(knn.pred, Direction.0910)
mean(knn.pred == Direction.0910)
```

```
      Direction.0910
knn.pred Down Up
   Down    17 21
   Up     26 40

0.548076923076923
```

```
In [14]: # KNN k = 100
set.seed(1)
knn.pred = knn(train.X, test.X, train.Direction, k=100)
table(knn.pred, Direction.0910)
mean(knn.pred == Direction.0910)
```

```
      Direction.0910
knn.pred Down Up
   Down    10 11
   Up     33 50

0.576923076923077
```

**Answer:** Out of these permutations, the original LDA and logistic regression have better performance in terms of test error rate.

## Exercise 12

This problem involves writing functions.

**12. a) Write a function, Power(), that prints out the result of raising 2 to the 3rd power. In other words, your function should compute  $2^3$  and print out the results.**

Hint: Recall that  $x^a$  raises  $x$  to the power  $a$ . Use the `print()` function to output the result.

```
In [15]: Power = function() {  
          2^3  
        }  
        print(Power())
```

```
[1] 8
```

**12 .b) Create a new function, Power2(), that allows you to pass any two numbers,  $x$  and  $a$ , and prints out the value of  $x^a$ . You can do this by beginning your function with the line ( `Power2 =function (x,a){` ) You should be able to call your function by entering, for instance, ( `Power2 (3,8)` ) on the command line. This should output the value of 38, namely, 6,561.**

```
In [16]: Power2 = function(x, a) {  
          x^a  
        }  
        Power2(3,8)
```

```
6561
```

**12. c) Using the Power2() function that you just wrote, compute  $10^3$ ,  $8^{17}$ , and  $131^3$ .**

```
In [17]: Power2(10, 3)  
        Power2(8, 17)  
        Power2(131, 3)
```

```
1000
```

```
2251799813685248
```


```
2248091
```

12. d) Now create a new function, `Power3()`, that actually returns the result  $x^a$  as an R object, rather than simply printing it to the screen. That is, if you store the value  $x^a$  in an object called `result` within your function, then you can simply `return()` this result, using the following line: `( return(result) )`. This should be the last line in your function, before the `}` symbol.

```
In [18]: Power3 = function(x, a) {  
  result = x^a  
  return(result)  
}
```

12. e) Now using the `Power3()` function, create a plot of  $f(x) = x^2$ . The x-axis should display a range of integers from 1 to 10, and the y-axis should display  $x^2$ . Label the axes appropriately, and use an appropriate title for the figure. Consider displaying either the x-axis, the y-axis, or both on the log-scale. You can do this by using `log="x"`, `log="y"`, or `log="xy"` as arguments to the `plot()` function.


```
In [19]: x = 1:10
par(mfrow=c(1,3))# graph parameters
plot(x, Power3(x, 2), log="x", ylab=expression(x^2), xlab=expression(x),
     main="log-transformed\nx-axis")
plot(x, Power3(x, 2), log="y", ylab=expression(x^2), xlab=expression(x),
     main="log-transformed\ny-axis")
plot(x, Power3(x, 2), log="xy", ylab=expression(x^2), xlab=expression(x),
     main="log-transformed\nx and y-axis")
par(mfrow=c(1,1))# reset graphic parameters
```



12. f) Create a function, `PlotPower()`, that allows you to create a plot of  $x$  against  $x^a$  for a fixed  $a$  and for a range of values of  $x$ . For instance, if you call `( PlotPower (1:10 ,3) )` then a plot should be created with an x-axis taking on values 1, 2, ..., 10, and a y-axis taking on values  $1^3, 2^3, \dots, 10^3$



```
In [20]: PlotPower = function(x, a) {  
  ylab_text <- bquote('x'^(a)) # write y-axis title  
  plot(x, Power3(x, a),  
        ylab = ylab_text)  
}  
PlotPower(1:10, 3)
```



## Exercise 13

Using the Boston data set, fit classification models in order to predict whether a given suburb has a crime rate above or below the median. Explore logistic regression, LDA, and KNN models using various subsets of the predictors. Describe your findings.

```
In [21]: summary(Boston)
attach(Boston)
crime01 = rep(0, length(crim))
crime01[crim>median(crim)] = 1
Boston = data.frame(Boston, crime01)

train = 1:(dim(Boston)[1]/2)
test = (dim(Boston)[1]/2+1):dim(Boston)[1]
Boston.train = Boston[train,]
Boston.test = Boston[test,]
crime01.test = crime01[test]
```

crim	zn	indus	chas
Min. : 0.00632	Min. : 0.00	Min. : 0.46	Min. : 0.00000
1st Qu.: 0.08204	1st Qu.: 0.00	1st Qu.: 5.19	1st Qu.: 0.00000
Median : 0.25651	Median : 0.00	Median : 9.69	Median : 0.00000
Mean : 3.61352	Mean : 11.36	Mean : 11.14	Mean : 0.06917
3rd Qu.: 3.67708	3rd Qu.: 12.50	3rd Qu.: 18.10	3rd Qu.: 0.00000
Max. : 88.97620	Max. : 100.00	Max. : 27.74	Max. : 1.00000

nox	rm	age	dis
Min. : 0.3850	Min. : 3.561	Min. : 2.90	Min. : 1.130
1st Qu.: 0.4490	1st Qu.: 5.886	1st Qu.: 45.02	1st Qu.: 2.100
Median : 0.5380	Median : 6.208	Median : 77.50	Median : 3.207
Mean : 0.5547	Mean : 6.285	Mean : 68.57	Mean : 3.795
3rd Qu.: 0.6240	3rd Qu.: 6.623	3rd Qu.: 94.08	3rd Qu.: 5.188
Max. : 0.8710	Max. : 8.780	Max. : 100.00	Max. : 12.127

rad	tax	ptratio	black
Min. : 1.000	Min. : 187.0	Min. : 12.60	Min. : 0.32
1st Qu.: 4.000	1st Qu.: 279.0	1st Qu.: 17.40	1st Qu.: 375.38
Median : 5.000	Median : 330.0	Median : 19.05	Median : 391.44
Mean : 9.549	Mean : 408.2	Mean : 18.46	Mean : 356.67
3rd Qu.: 24.000	3rd Qu.: 666.0	3rd Qu.: 20.20	3rd Qu.: 396.23
Max. : 24.000	Max. : 711.0	Max. : 22.00	Max. : 396.90

lstat	medv
Min. : 1.73	Min. : 5.00
1st Qu.: 6.95	1st Qu.: 17.02
Median : 11.36	Median : 21.20
Mean : 12.65	Mean : 22.53
3rd Qu.: 16.95	3rd Qu.: 25.00
Max. : 37.97	Max. : 50.00

```
In [22]: # logistic regression
glm.fit = glm(crime01~.-crime01-crim,
              data=Boston, family=binomial, subset=train)
glm.probs = predict(glm.fit, Boston.test, type="response")
glm.pred = rep(0, length(glm.probs))
glm.pred[glm.probs > 0.5] = 1
mean(glm.pred != crime01.test)
```

0.181818181818182

**Answer:** 18.2% test error rate.

```
In [23]: glm.fit = glm(crime01~.-crime01-crim-chas-tax,
                      data=Boston, family=binomial, subset=train)
glm.probs = predict(glm.fit, Boston.test, type="response")
glm.pred = rep(0, length(glm.probs))
glm.pred[glm.probs > 0.5] = 1
mean(glm.pred != crime01.test)
```

0.185770750988142

**Answer:** 18.6% test error rate.

```
In [24]: # LDA
lda.fit = lda(crime01~.-crime01-crim, data=Boston, subset=train)
lda.pred = predict(lda.fit, Boston.test)
mean(lda.pred$class != crime01.test)
```

0.134387351778656

**Answer:** 13.4% test error rate.

```
In [25]: lda.fit = lda(crime01~.-crime01-crim-chas-tax, data=Boston, subset=train)
lda.pred = predict(lda.fit, Boston.test)
mean(lda.pred$class != crime01.test)
```

0.122529644268775

**Answer:** 12.3% test error rate.

```
In [26]: lda.fit = lda(crime01~.-crime01-crim-chas-tax-lstat-indus-age,
                        data=Boston, subset=train)
lda.pred = predict(lda.fit, Boston.test)
mean(lda.pred$class != crime01.test)
```

0.118577075098814

**Answer:** 11.9% test error rate.

```
In [27]: # KNN
library(class)
train.X = cbind(zn, indus, chas, nox, rm, age, dis, rad, tax, ptratio, black,
                lstat, medv)[train,]
test.X = cbind(zn, indus, chas, nox, rm, age, dis, rad, tax, ptratio, black,
               lstat, medv)[test,]
train.crime01 = crime01[train]
set.seed(1)
# KNN(k=1)
knn.pred = knn(train.X, test.X, train.crime01, k=1)
mean(knn.pred != crime01.test)
```

0.458498023715415

**Answer:** 45.8% test error rate.

```
In [28]: # KNN(k=10)
set.seed(1)
knn.pred = knn(train.X, test.X, train.crime01, k=10)
mean(knn.pred != crime01.test)
```

0.110671936758893

**Answer:** 11.1% test error rate.

```
In [29]: # KNN(k=100)
set.seed(1)
knn.pred = knn(train.X, test.X, train.crime01, k=100)
mean(knn.pred != crime01.test)
```

0.486166007905138

**Answer:** 48.6% test error rate.

In general, the best models are the ones with the smaller test error rates. In our case, this means that the lda.fit and the KNN with K=10 are the best models.